

# **CARTE ML Workshop**


Lecture 2-1: Unsupervised Learning and Visualization

# Clustering


The most common type of unsupervised learning

Goal: group “similar” data points together

Unsupervised because we don't label the data as we did in classification/regression: let the features speak for themselves!




Clustering galaxies, from Miller et al. (2005)




# Data Preparation for Clustering

The points are colored by a  
clustering algorithm

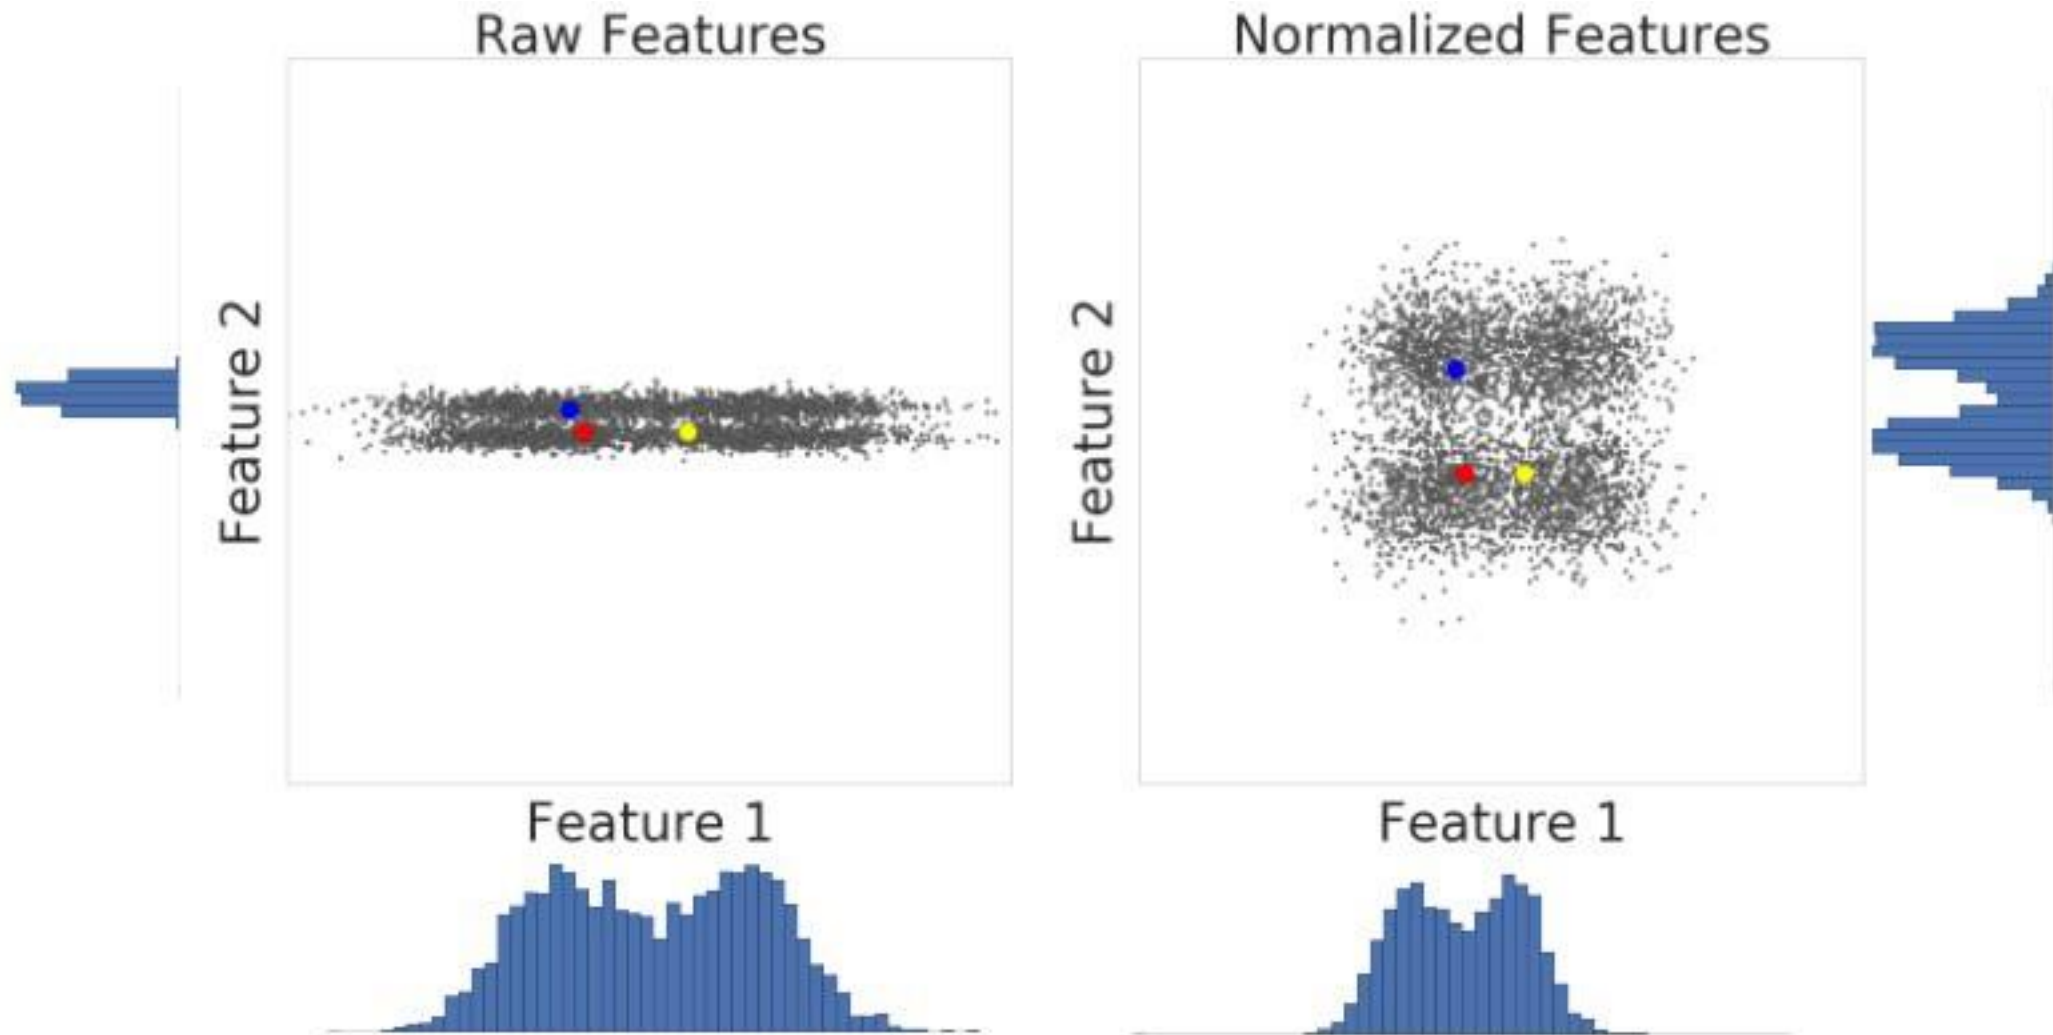
Raw data



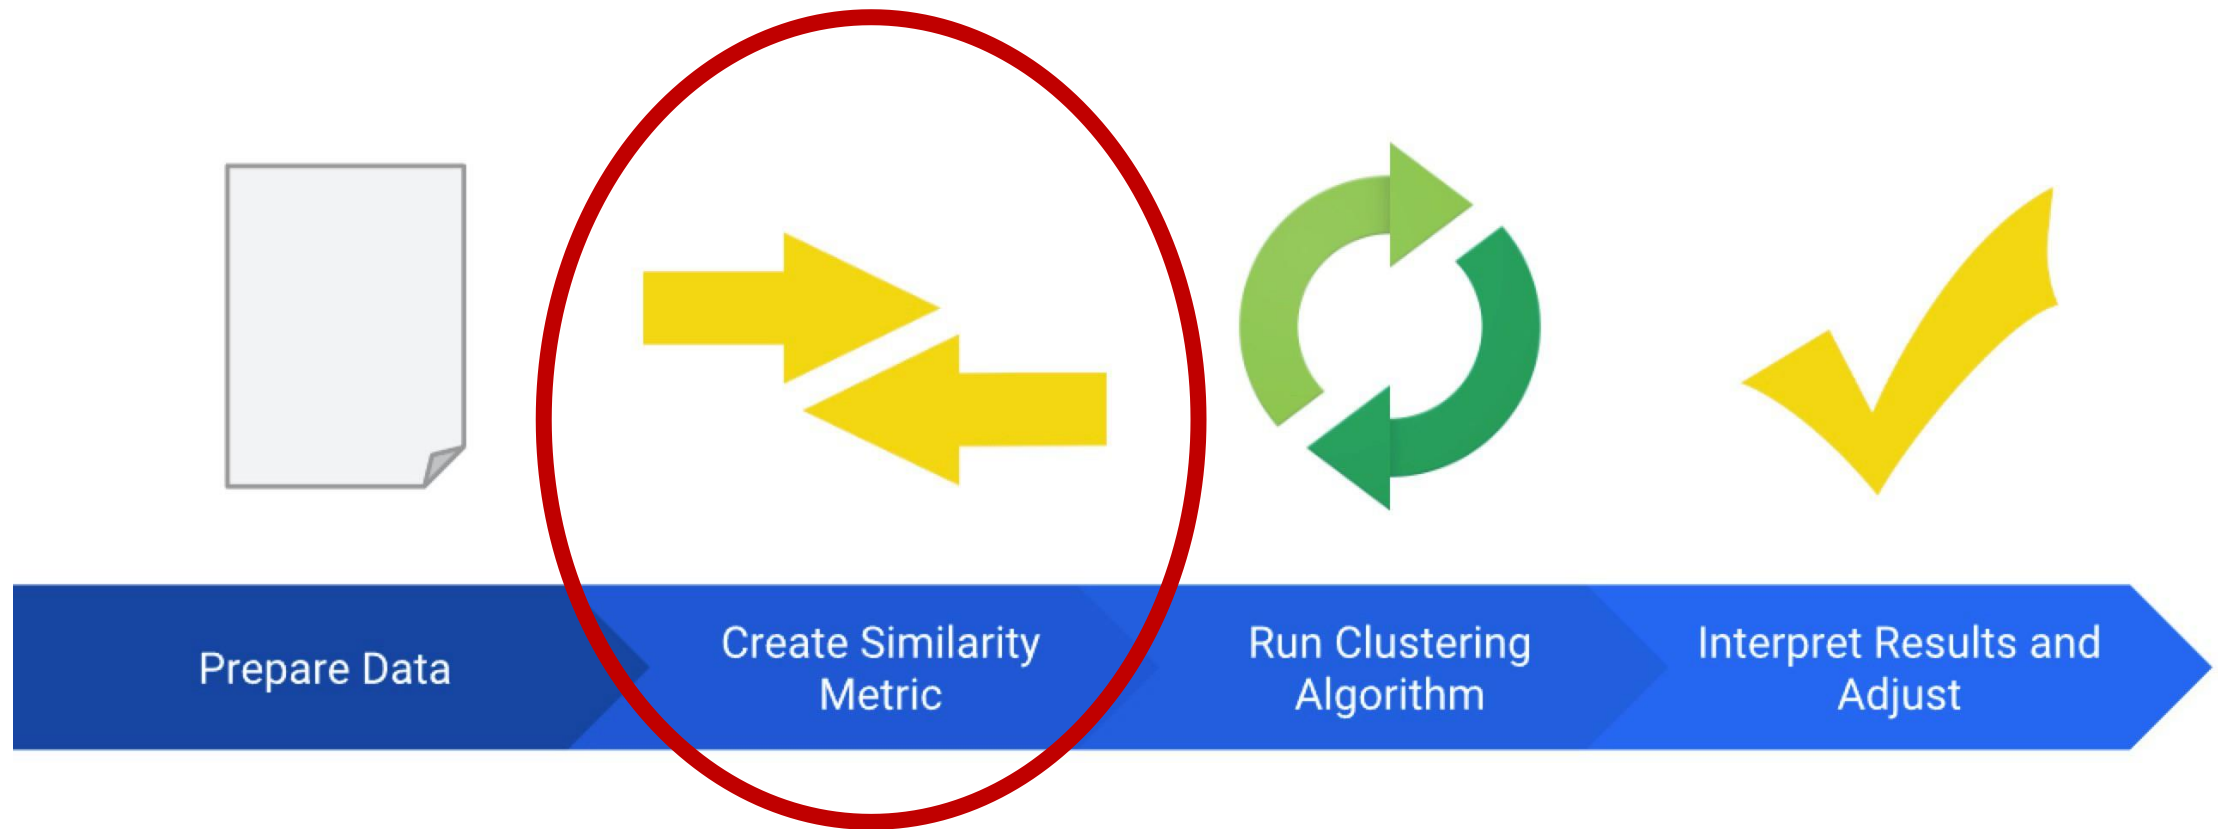
Standardized data



# Data Preparation for Clustering



# Clustering workflow



From Google's Clustering lesson: <https://developers.google.com/machine-learning/clustering/>

# Distance Metrics

The right distance metric depends on your application!

Distance of vectors  $x = (x_1, \dots, x_n)$  and  $y = (y_1, \dots, y_n)$

- Euclidean distance

$$d(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

- Manhattan distance

$$d(x, y) = \sum_{i=1}^n |x_i - y_i|$$

- Correlation distance

$$d(x, y) = 1 - r(x, y)$$

$r(x, y)$  is Pearson correlation coefficient

Distance of sequences ACCTTG and TACCTG

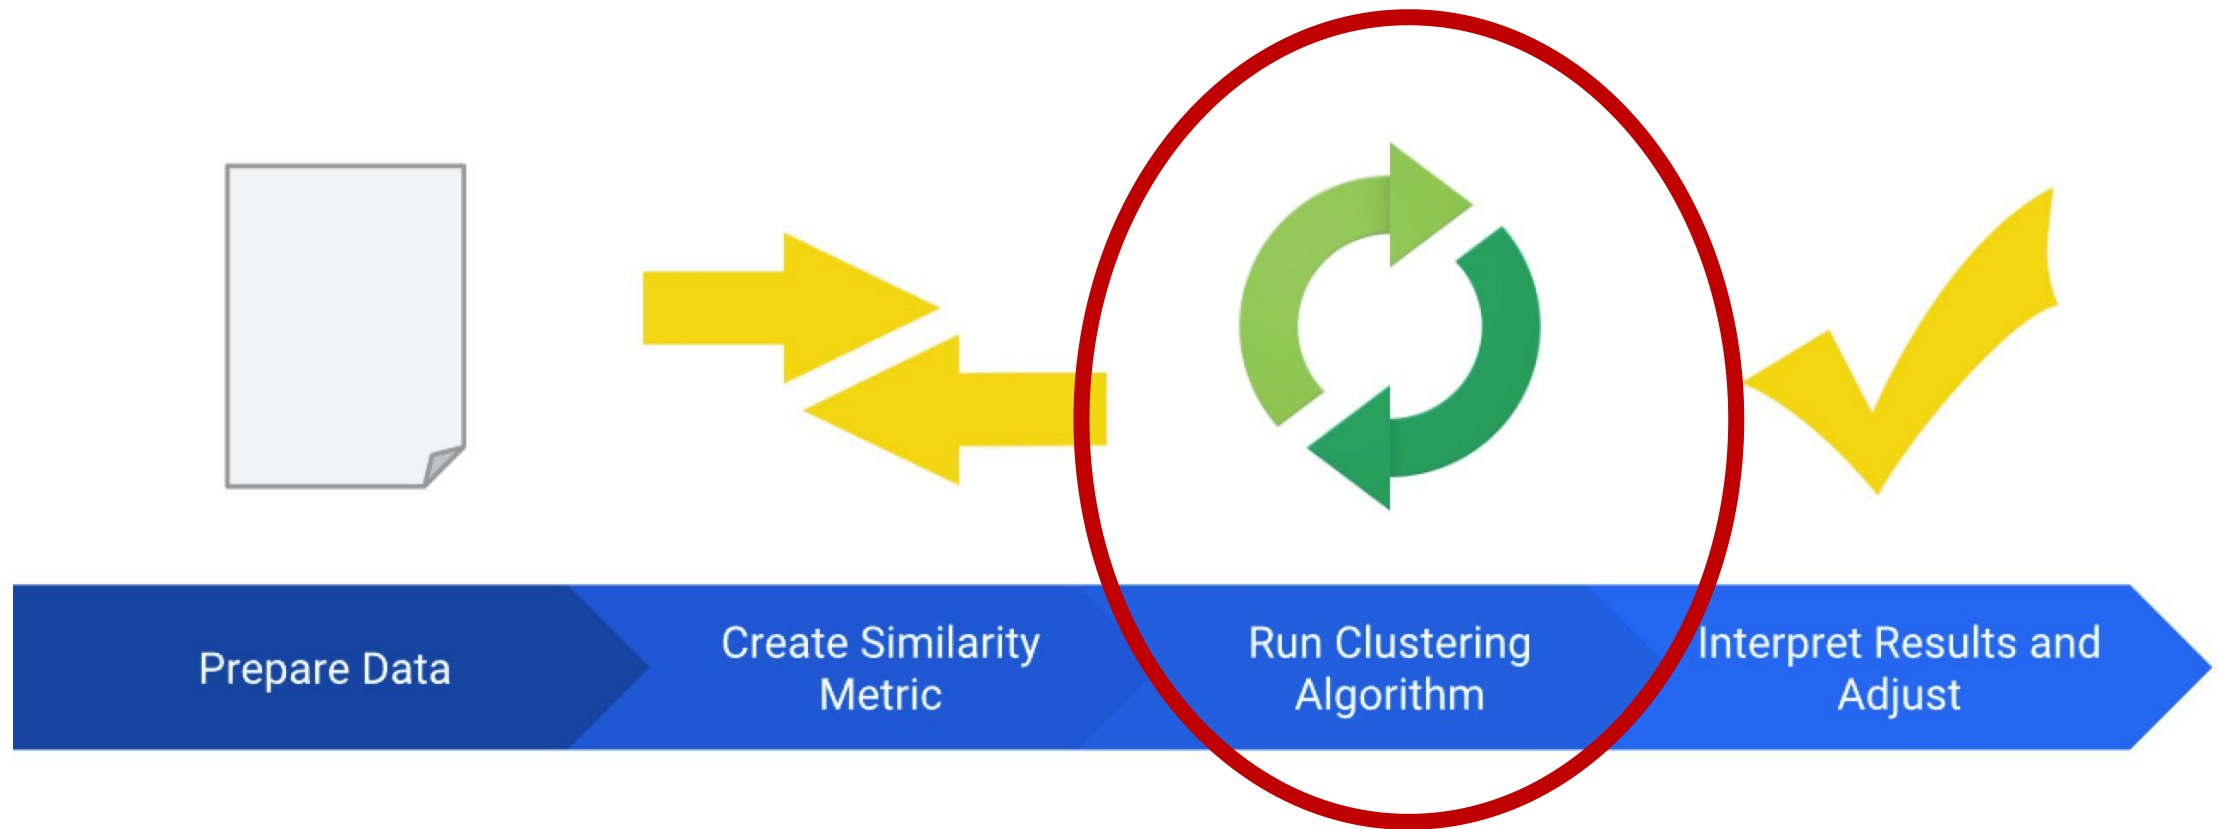
- Hamming distance

$$\begin{array}{c} \underline{\mathbf{A}}\mathbf{C}\mathbf{C}\mathbf{T}\underline{\mathbf{T}}\mathbf{G} \\ \underline{\mathbf{T}}\mathbf{A}\mathbf{C}\mathbf{C}\underline{\mathbf{T}}\mathbf{G} \end{array} \Rightarrow 3$$

- Levenshtein distance

$$\begin{array}{c} \underline{\mathbf{.}}\mathbf{A}\mathbf{C}\mathbf{C}\mathbf{T}\underline{\mathbf{T}}\mathbf{G} \\ \underline{\mathbf{T}}\mathbf{A}\mathbf{C}\mathbf{C}\underline{\mathbf{.}}\mathbf{T}\mathbf{G} \end{array} \Rightarrow 2$$

# Clustering workflow




From Google's Clustering lesson: <https://developers.google.com/machine-learning/clustering/>

# K-means Clustering

Euclidean distance metric

Given a set of data points...






# K-means Clustering

Euclidean distance metric


Select  $k=3$  initial centers at random



# K-means Clustering

Euclidean distance metric


Recompute optimal centers given a fixed clustering



# K-means Clustering

Euclidean distance metric


Assign each point to its nearest center



# K-means Clustering

Euclidean distance metric


Recompute optimal centers given a fixed clustering



# K-means Clustering

Euclidean distance metric


Assign each point to its nearest center




# K-means Clustering

Euclidean distance metric

Recompute optimal centers given a fixed clustering




# Selecting the number of clusters




**FIGURE 14.8.** Total within-cluster sum of squares for  $K$ -means clustering applied to the human tumor microarray data.

# K-means can fail



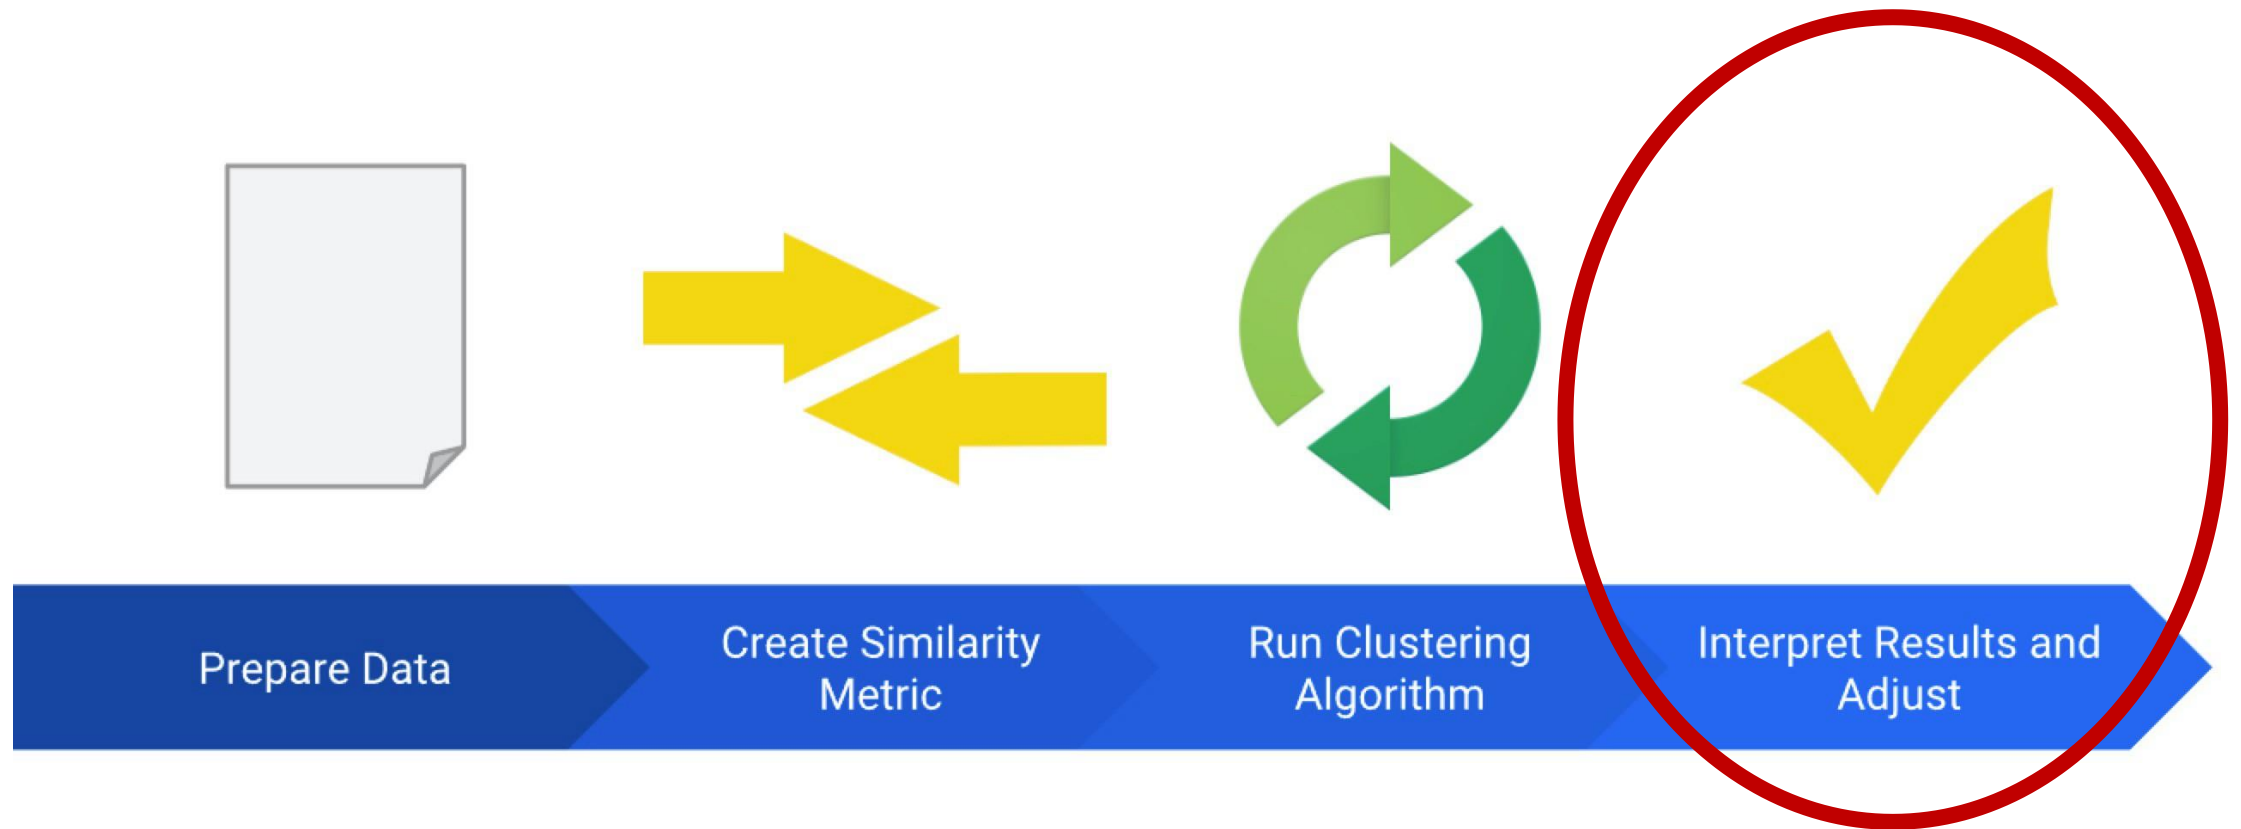
This is a heuristic!  
No guarantees it'll find optimum



In practice, smarter centroid  
initialization solves this





# Clustering workflow



From Google's Clustering lesson: <https://developers.google.com/machine-learning/clustering/>

# Hierarchical Clustering

High-level idea: build a tree (hierarchy) of clusters




# Hierarchical Clustering

# Human Tumor Microarray Data

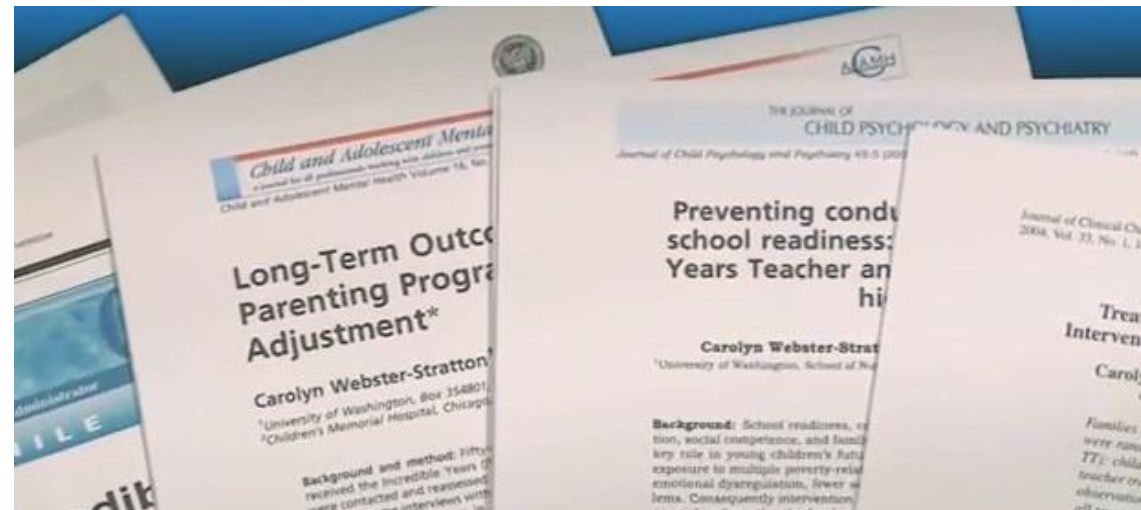


64 samples, 6830 features  
(gene expression levels)

# Elements of Statistical Learning, Chapter 14.3.8




# Dimensionality Reduction

or dealing with very high-dimensional data



# Dimensionality Reduction

Examples: Principal Component Analysis (PCA), t-SNE, ...



Powerful unsupervised learning techniques for extracting hidden (potentially lower dimensional) structure from high dimensional datasets.

Useful for:


- Visualization
- Data compression for faster supervised learning
- Noise removal

Based on slide by Nina Balcan



# Principal Component Analysis (PCA)

- Principal Components (PC) are orthogonal directions that capture most of the variance in the data.






# PCA for Face Reconstruction

Eigenfaces, slide based on Derek Hoiem's, UIUC CS543

Image dataset



64 Principal Components



PCA algorithm



# PCA for Face Reconstruction

Eigenfaces, slide based on Derek Hoesim's, UIUC CS543

Face Reconstruction using the Principal Components



=



+





# PCA for Face Reconstruction

Eigenfaces, slide based on Derek Hoesim's, UIUC CS543

Face Recognition using PCA:

- 1 Given face image datasets, extract Principal Components  $v_1, v_2, \dots, v_n$ .
- 2 Given new image, project onto PCs.
- 3 Find closest (projected) image in training dataset



[\[PDF\] Face recognition using eigenfaces](#)



[MA Turk, AP Pentland](#) - ... on Computer Vision and Pattern **Recognition**, 1991 - cin.ufpe.br

We present an approach to the detection and identification of human faces and describe a working, near-real-time face recognition system which tracks a subject's head and then recognizes the person by comparing characteristics of the face to those of known ...

☆ 🔖 Cited by 7225 Related articles All 61 versions 🔗

# Final words on PCA

- Advantages
  - Fast to compute an optimal solution: an eigenvector problem
  - No hyper-parameters to tune
- Caveats
  - Discards information
  - Limited to linear projections



- From Michael Guerzhoy's slides, UofT CSC320


# t-SNE

t-Distributed Stochastic Neighbor Embedding


784 dimensions = 28x28 pixels

2 dimensions

Random Sampling of MNIST




t-SNE



# t-SNE


- Create a distribution representing the similarity of points
- Similarity of A to B is the conditional probability  $p_{A|B}$  that B would pick A as its neighbour



<https://towardsdatascience.com/t-sne-clearly-explained-d84c537f53a>

# t-SNE

- Create a matching distribution in the low dimensional space, such that the divergence between  $p_{A|B}$  and  $q_{A|B}$  is minimized
- Key hyperparameter: width of the gaussian is controlled by *perplexity*




<https://towardsdatascience.com/t-sne-clearly-explained-d84c537f53a>

# t-SNE

- Last step: while we use a **Gaussian** in the original space, we use a **Student's t-distribution** in the lower dimensional space
- Intuition: because we are losing dimensions, we need more distances that map to approximately the same probability
- Student's t-distribution has a longer tail than a Gaussian, allowing for more “options” in how the result is organized

# UMAP

- Both UMAP and t-SNE try to find a low-dimensional representation that matches the structure of the high-dimensional one
- But while t-SNE looks at pairwise distances between points, UMAP looks at a network of neighbours





# More info

- <https://pair-code.github.io/understanding-umap/>